

Class 10: Halloween Mini-Project

Rosa Chavez (PID: A16487039)

```
options(repos = c(CRAN="http://cran.us.r-project.org"))
```

```
candy_file <- "https://raw.githubusercontent.com/fivethirtyeight/data/master/candy-power-r
```

```
candy <- read.csv(candy_file, row.names = 1)
head(candy)
```

	chocolate	fruity	caramel	peanutyalmondy	nougat	crispedricewafer
100 Grand	1	0	1	0	0	1
3 Musketeers	1	0	0	0	1	0
One dime	0	0	0	0	0	0
One quarter	0	0	0	0	0	0
Air Heads	0	1	0	0	0	0
Almond Joy	1	0	0	1	0	0

	hard	bar	pluribus	sugarpercent	pricepercent	winpercent
100 Grand	0	1	0	0.732	0.860	66.97173
3 Musketeers	0	1	0	0.604	0.511	67.60294
One dime	0	0	0	0.011	0.116	32.26109
One quarter	0	0	0	0.011	0.511	46.11650
Air Heads	0	0	0	0.906	0.511	52.34146
Almond Joy	0	1	0	0.465	0.767	50.34755

Q1. How many different candy types are in this dataset?

```
n_candy_types <- nrow(candy)
n_candy_types
```

```
[1] 85
```

Q2. How many fruity candy types are in the dataset?

```
n_fruity_candies <- sum(candy$fruity == 1)
n_fruity_candies
```

```
[1] 38
```

Q3. What is your favorite candy in the dataset and what is its winpercent value?

Hershey's Special Dark

```
winpercent_fav <- candy["Hershey's Special Dark", "winpercent"]
winpercent_fav
```

```
[1] 59.23612
```

Q4. What is the winpercent value for "Kit Kat"?

```
winpercent_kit_kat <- candy["Kit Kat", "winpercent"]
winpercent_kit_kat
```

```
[1] 76.7686
```

Q5. What is the winpercent value for "Tootsie Roll Snack Bars"?

```
winpercent_tootsie_roll <- candy["Tootsie Roll Snack Bars", "winpercent"]
winpercent_tootsie_roll
```

```
[1] 49.6535
```

```
install.packages("skimr")
```

Installing package into 'C:/Users/ITSloaner.DESKTOP-I89K3M9/AppData/Local/R/win-library/4.3'
(as 'lib' is unspecified)

package 'skimr' successfully unpacked and MD5 sums checked

The downloaded binary packages are in

C:\Users\ITSloaner.DESKTOP-I89K3M9\AppData\Local\Temp\RtmpQjtUVP\downloaded_packages

```
library("skimr")
```

```
skim(candy)
```

Table 1: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12
Group variables	None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmondy	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?


Yes, the last one winpercent does not have numbers 0 to 1 but rather 0 to 100.

Q7. What do you think a zero and one represent for the candy\$chocolate column?

1 represents candies that contain chocolates while 0 represents candies that do not contain chocolate.

Q8. Plot a histogram of winpercent values

```
hist(candy$winpercent, main="Histogram of Winpercent", xlab="Win Percent")
```



Q9. Is the distribution of winpercent values symmetrical?

The distribution is slightly skewed to the right.

Q10. Is the center of the distribution above or below 50%?

Center is above 50%.

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

```
avg_chocolate <- mean(candy$winpercent[candy$chocolate == 1])
avg_fruity <- mean(candy$winpercent[candy$fruity == 1])
avg_chocolate
```

```
[1] 60.92153
```

```
avg_fruity
```

```
[1] 44.11974
```

chocolate is higher ranked than fruity candy.

Q12. Is this difference statistically significant?

```
t.test(candy$winpercent[candy$chocolate == 1], candy$winpercent[candy$fruity == 1])
```

Welch Two Sample t-test

```
data: candy$winpercent[candy$chocolate == 1] and candy$winpercent[candy$fruity == 1]
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
 60.92153  44.11974
```

Q13. What are the five least liked candy types in this set?

```
least_liked <- head(candy[order(candy$winpercent), ], n = 5)
least_liked
```

	chocolate	fruity	caramel	peanut	almond	nougat
Nik L Nip	0	1	0		0	0
Boston Baked Beans	0	0	0		1	0
Chiclets	0	1	0		0	0
Super Bubble	0	1	0		0	0
Jawbusters	0	1	0		0	0

	crisped	rice	wafer	hard	bar	pluribus	sugar	percent	price	percent
Nik L Nip		0	0	0		1		0.197		0.976
Boston Baked Beans		0	0	0		1		0.313		0.511
Chiclets		0	0	0		1		0.046		0.325
Super Bubble		0	0	0		0		0.162		0.116
Jawbusters		0	1	0		1		0.093		0.511

	winpercent
Nik L Nip	22.44534
Boston Baked Beans	23.41782
Chiclets	24.52499

Super Bubble	27.30386
Jawbusters	28.12744

Q14. What are the top 5 all time favorite candy types out of this set?

```
most_liked <- head(candy[order(candy$winpercent, decreasing = TRUE), ], n = 5)
most_liked
```


	chocolate	fruity	caramel	peanut	almond	nougat
Reese's Peanut Butter cup	1	0	0		1	0
Reese's Miniatures	1	0	0		1	0
Twix	1	0	1		0	0
Kit Kat	1	0	0		0	0
Snickers	1	0	1		1	1

	crisped	rice	wafer	hard bar	pluribus	sugar	percent
Reese's Peanut Butter cup		0	0	0	0		0.720
Reese's Miniatures		0	0	0	0		0.034
Twix		1	0	1	0		0.546
Kit Kat		1	0	1	0		0.313
Snickers		0	0	1	0		0.546

	price	percent	winpercent
Reese's Peanut Butter cup	0.651		84.18029
Reese's Miniatures	0.279		81.86626
Twix	0.906		81.64291
Kit Kat	0.511		76.76860
Snickers	0.651		76.67378


Q15. Make a first barplot of candy ranking based on winpercent values.

```
library(ggplot2)
ggplot(candy) +
  aes(winpercent, rownames(candy)) +
  geom_bar(stat = "identity")
```



Q16. This is quite ugly, use the `reorder()` function to get the bars sorted by winpercent?

```
library(ggplot2)
ggplot(candy, aes(x = reorder(rownames(candy), -winpercent), y = winpercent)) +
  geom_bar(stat = "identity") +
  coord_flip() +
  theme_minimal()
```



Q17. What is the worst ranked chocolate candy?

```
worst_chocolate <- candy[candy$chocolate == 1, ] [which.min(candy[candy$chocolate == 1, "winpercent"]), ]
worst_chocolate
```

```
      chocolate fruity caramel peanutyalmondy nougat crispedricewafer hard
Sixlets         1         0         0         0         0         0         0
      bar pluribus sugarpercent pricepercent winpercent
Sixlets         0         1         0.22         0.081         34.722
```

Q18. What is the best ranked fruity candy?

```
best_fruity <- candy[candy$fruity == 1, ] [which.max(candy[candy$fruity == 1, "winpercent"]), ]
best_fruity
```

```
      chocolate fruity caramel peanutyalmondy nougat crispedricewafer hard
Starburst         0         1         0         0         0         0         0
      bar pluribus sugarpercent pricepercent winpercent
Starburst         0         1         0.151         0.22         67.03763
```



```
install.packages("ggrepel")
```

Installing package into 'C:/Users/ITSloaner.DESKTOP-I89K3M9/AppData/Local/R/win-library/4.3'
(as 'lib' is unspecified)

package 'ggrepel' successfully unpacked and MD5 sums checked

The downloaded binary packages are in

C:\Users\ITSloaner.DESKTOP-I89K3M9\AppData\Local\Temp\RtmpQjtUVP\downloaded_packages


```
library(ggrepel)
```

Q19. Which candy type is the highest ranked in terms of winpercent for the least money -
i.e. offers the most bang for your buck?


```
my_cols=rep("black", nrow(candy))
my_cols[as.logical(candy$chocolate)] = "chocolate"
my_cols[as.logical(candy$bar)] = "brown"
my_cols[as.logical(candy$fruity)] = "pink"
```

```
ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, size=3.3, max.overlaps = 5)
```

Warning: ggrepel: 65 unlabeled data points (too many overlaps). Consider
increasing max.overlaps




```
ggplot(candy) +
  aes(winpercent, reorder(rownames(candy),winpercent)) +
  geom_col(fill=my_cols)
```



```
ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, size=3.3, max.overlaps = 5)
```

Warning: ggrepel: 65 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Reese'e Minatures

Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```
ord <- order(candy$pricepercent, decreasing = TRUE)
expensive_candies <- head(candy[ord, ], n = 5)
print(expensive_candies)
```

	chocolate	fruity	caramel	peanutyalmondy	nougat
Nik L Nip	0	1	0	0	0
Nestle Smarties	1	0	0	0	0
Ring pop	0	1	0	0	0
Hershey's Krackel	1	0	0	0	0
Hershey's Milk Chocolate	1	0	0	0	0

	crispedricewafer	hard	bar	pluribus	sugarpercent
Nik L Nip	0	0	0	1	0.197
Nestle Smarties	0	0	0	1	0.267
Ring pop	0	1	0	0	0.732
Hershey's Krackel	1	0	1	0	0.430
Hershey's Milk Chocolate	0	0	1	0	0.430

pricepercent winpercent

Nik L Nip	0.976	22.44534
Nestle Smarties	0.976	37.88719
Ring pop	0.965	35.29076
Hershey's Krackel	0.918	62.28448
Hershey's Milk Chocolate	0.918	56.49050

```
install.packages("corrplot")
```

Installing package into 'C:/Users/ITSloaner.DESKTOP-I89K3M9/AppData/Local/R/win-library/4.3'
(as 'lib' is unspecified)

package 'corrplot' successfully unpacked and MD5 sums checked


The downloaded binary packages are in

C:\Users\ITSloaner.DESKTOP-I89K3M9\AppData\Local\Temp\RtmpQjtUVP\downloaded_packages

```
library(corrplot)
```

corrplot 0.92 loaded

```
cij <- cor(candy)  
corrplot(cij)
```



Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

Chocolate and sugarpercent.

Q23. Similarly, what two variables are most positively correlated?

Chocolate and nougat.

Q24. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you?


```
pca <- prcomp(candy[, -1], scale. = TRUE)
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	1.9200	1.1143	1.1085	1.0751	0.95010	0.81815	0.81352
Proportion of Variance	0.3351	0.1129	0.1117	0.1051	0.08206	0.06085	0.06016
Cumulative Proportion	0.3351	0.4480	0.5597	0.6648	0.74685	0.80770	0.86787

	PC8	PC9	PC10	PC11
Standard deviation	0.68950	0.64410	0.60875	0.43887
Proportion of Variance	0.04322	0.03772	0.03369	0.01751
Cumulative Proportion	0.91109	0.94880	0.98249	1.00000

```
plot(pca$x[, 1:2], col = my_cols)
```



```
# Adding color and labels to PCA plot
ggplot(as.data.frame(pca$x[, 1:2]), aes(x = PC1, y = PC2, color = my_cols, label = rowname)) +
  geom_point() +
  geom_text_repel()
```

Warning: ggrepel: 54 unlabeled data points (too many overlaps). Consider increasing max.overlaps

